

# A PINN-Based Framework for Efficient European Option Pricing under the Black-Scholes Model

Yitian Liang<sup>1\*</sup>

<sup>1\*</sup>Department of Applied Physics and Applied Mathematics, Columbia University, New York, 10027, NY, USA.

Corresponding author(s). E-mail(s): [y15433@columbia.edu](mailto:y15433@columbia.edu);

## Abstract

This paper introduces a physics-informed neural network (PINN) approach for solving the Black-Scholes partial differential equation (PDE) associated with European call option pricing. By directly integrating the PDE and its terminal and boundary conditions into the loss function, the PINN framework achieves an accurate, grid-free approximation of the option price surface. Benchmarking against the known closed-form European call solution verifies that the trained PINN converges smoothly to the theoretical result, capturing the correct payoff at maturity and maintaining stability across the entire domain. The resulting method provides a flexible and efficient alternative to traditional numerical schemes, reducing complexity while retaining high accuracy and reliability in a range of parameter settings.

**Keywords:** PINNs, Black-Scholes PDE, European call option, Option pricing, Physics-informed machine learning

## 1 Introduction

Accurate and efficient option pricing is a fundamental goal in quantitative finance. Traditional approaches often rely on numerical methods, such as finite differences or Monte Carlo simulations, which can become computationally expensive or challenging to scale. Although the classical Black-Scholes equation admits a closed-form solution for European options, this advantage does not extend to more complex instruments or parameter variations. As markets evolve and novel payoff structures emerge, there is a need for more flexible and data-driven techniques that can incorporate both theoretical insights and practical constraints.

Physics-informed neural networks (PINNs) have recently been proposed as a promising alternative. Unlike standard deep learning methods that primarily fit observed data, PINNs incorporate the governing partial differential equations (PDEs) and associated conditions into the loss function. This ensures that the training process is guided by fundamental principles rather than relying on purely data-driven patterns. Applications in finance have demonstrated the potential of PINNs, including their ability to handle complex fractional Black-Scholes equations and free boundary problems [1], as well as standard American and European option pricing scenarios [2].

The present work focuses on applying a PINN-based approach to the classical European call option pricing problem. By embedding the Black-Scholes PDE, along with its terminal and boundary conditions, directly into the learning process, the PINN avoids traditional discretization schemes. The resulting solution converges to the known closed-form formula at maturity, ensuring both accuracy and interpretability. In addition, the flexible nature of PINNs allows for adjustments to parameter settings without extensive remodeling, offering a powerful framework for future extensions and more complex derivatives.

The remainder of this paper is organized as follows: Section 2 reviews the Black-Scholes model and the European call option's closed-form solution. Section 3 outlines the PINN approach, detailing how the PDE and conditions are integrated into the training procedure. Section 4 presents numerical experiments, benchmarking the PINN results against analytical solutions and discussing the outcomes. Finally, Section 5 offers a broader reflection on these findings and their potential implications.

## 2 Background and Theoretical Foundations

The Black-Scholes framework provides a well-established theoretical foundation for pricing European call options under idealized market assumptions. In this model, the option value  $V(S, t)$ , with  $S$  as the spot price and  $t$  as the current time, evolves according to the well-known Black-Scholes partial differential equation (PDE):

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} - rV = 0, \quad (1)$$

where  $r$  is the risk-free interest rate and  $\sigma$  is the constant volatility. For a European call option with strike  $K$  and maturity  $T$ , the terminal condition at  $t = T$  is given by:

$$V(S, T) = \max(S - K, 0). \quad (2)$$

Under these assumptions, the Black-Scholes equation admits a closed-form solution:

$$C(S, t) = SN(d_1) - Ke^{-r(T-t)}N(d_2), \quad (3)$$

with

$$d_1 = \frac{\ln(S/K) + (r + \frac{\sigma^2}{2})(T - t)}{\sigma\sqrt{T - t}}, \quad d_2 = d_1 - \sigma\sqrt{T - t}, \quad (4)$$

and  $N(\cdot)$  being the standard normal cumulative distribution function. This closed-form solution serves as a benchmark against which numerical or approximate methods can be compared.

The ability to verify approximate solutions against the exact Black-Scholes formula is crucial. It provides a clear metric for evaluating new computational approaches, including the physics-informed neural network (PINN) method, ensuring that the network’s approximation aligns with well-established theoretical results.

## 3 Methodology

### 3.1 Physics-Informed Neural Networks (PINNs)

Physics-Informed Neural Networks (PINNs) combine deep neural network approximation capabilities with the intrinsic structure of partial differential equations (PDEs). Instead of merely fitting to data points, PINNs embed the governing PDE, along with associated boundary and initial (or terminal) conditions, directly into the loss function. This approach ensures that the network’s training process is guided by the underlying physics or financial mathematics rather than relying solely on observed data. For option pricing, PINNs provide a flexible framework that can incorporate the Black-Scholes PDE without the need for predefined grids or explicit time-stepping schemes. As a result, PINNs can efficiently solve high-dimensional problems and seamlessly integrate complex or time-varying parameters.

### 3.2 PINN for European Call Option Pricing

We focus on the European call option, a fundamental problem in quantitative finance. Let  $V(S, t)$  represent the option value, with  $S$  as the spot price and  $t$  as time. Under the Black-Scholes framework,  $V(S, t)$  satisfies the following PDE:

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} - rV = 0. \quad (5)$$

Here,  $\sigma$  is the volatility of the underlying asset and  $r$  is the risk-free interest rate. For a European call option with strike  $K$  and maturity  $T$ , the terminal payoff is given by  $V(S, T) = \max(S - K, 0)$ .

To approximate the solution, the PINN takes  $(S, t)$  as input and outputs the corresponding option price. During training, the network parameters are adjusted to minimize violations of the PDE, terminal conditions, and spatial boundary conditions.

### 3.3 Problem Setup and Data Sampling

We consider  $S \in [0, S_{\max}]$  and  $t \in [0, T]$ . The choice of a sufficiently large  $S_{\max}$  approximates the behavior as  $S \rightarrow \infty$ . For the European call, the prescribed conditions are: - At  $S = 0$ ,  $V(0, t) = 0$ . - At  $S = S_{\max}$ ,  $V(S_{\max}, t) \approx S_{\max} - Ke^{-r(T-t)}$ .

At maturity  $t = T$ , the terminal payoff is  $V(S, T) = \max(S - K, 0)$ . We sample a set of collocation points  $(S_i, t_i)$  throughout the domain to enforce the PDE residual and additional points along boundaries and at  $t = T$  to enforce terminal and boundary conditions. These sampled points ensure coverage of the solution space, helping

the PINN converge to a function that respects both the governing equation and the imposed constraints.

### 3.4 Loss Function and Training Procedure

The total loss function integrates contributions from the PDE residual, terminal condition, and boundary conditions. Let  $V_\theta(S, t)$  denote the neural network approximation of  $V(S, t)$ .

The PDE residual mean squared error (MSE) is:

$$\text{MSE}_{\text{pde}} = \frac{1}{N_{\text{pde}}} \sum_{i=1}^{N_{\text{pde}}} \left( \frac{\partial V_\theta}{\partial t}(S_i, t_i) + \frac{1}{2} \sigma^2 S_i^2 \frac{\partial^2 V_\theta}{\partial S^2}(S_i, t_i) + r S_i \frac{\partial V_\theta}{\partial S}(S_i, t_i) - r V_\theta(S_i, t_i) \right)^2. \quad (6)$$

The terminal condition MSE is:

$$\text{MSE}_{\text{term}} = \frac{1}{N_{\text{term}}} \sum_{j=1}^{N_{\text{term}}} (V_\theta(S_j, T) - \max(S_j - K, 0))^2. \quad (7)$$

The boundary condition MSE combines conditions at  $S = 0$  and  $S = S_{\text{max}}$ :

$$\text{MSE}_{\text{bvp}} = \text{MSE}_{S=0} + \text{MSE}_{S=S_{\text{max}}}. \quad (8)$$

Weighted by respective coefficients  $\lambda_{\text{pde}}$ ,  $\lambda_{\text{term}}$ , and  $\lambda_{\text{bvp}}$ , the total loss is:

$$\text{Loss} = \lambda_{\text{pde}} \text{MSE}_{\text{pde}} + \lambda_{\text{term}} \text{MSE}_{\text{term}} + \lambda_{\text{bvp}} \text{MSE}_{\text{bvp}}. \quad (9)$$

We employ gradient-based optimization (e.g., Adam) to iteratively update the network parameters. By backpropagating through the PDE-constrained loss, the PINN converges toward a solution that satisfies both the theoretical financial model and the imposed conditions.

### 3.5 Benchmarking with the Analytical Black-Scholes Solution

To validate the accuracy of our PINN approach, we compare its predictions against the closed-form Black-Scholes solution for a European call option:

$$C(S, t) = S N(d_1) - K e^{-r(T-t)} N(d_2), \quad (10)$$

where

$$d_1 = \frac{\ln(S/K) + (r + \frac{\sigma^2}{2})(T-t)}{\sigma \sqrt{T-t}}, \quad d_2 = d_1 - \sigma \sqrt{T-t}. \quad (11)$$

Here,  $N(\cdot)$  is the standard normal cumulative distribution function. The closed-form solution serves as an exact benchmark. By evaluating the differences between the PINN output and this analytic reference, we quantify approximation errors and assess convergence.

In addition, we present various visualizations: - Time series comparisons of SPX spot and call option prices to contextualize the model predictions. - Illustrations of the sampling scheme showing how the collocation, boundary, and terminal points are chosen. - Training curve plots depicting loss evolution over epochs. - Three-dimensional surfaces and two-dimensional cross sections of PINN predictions, highlighting their consistency with the analytical solution.

These benchmarks and visual analyses confirm the PINN’s reliability, paving the way for extensions to more complex derivatives and market conditions.

## 4 Experiments and Results

In this section, we present and analyze the outcomes obtained by applying the PINN-based approach to the European call option pricing problem. All results are directly derived from the code and setup described in the previous sections, ensuring consistency between the theoretical foundations, methodology, and empirical validations.

**Table 1** Simulation Data Parameters


Parameter	Value
Strike Price, $K$	40
Risk Free Rate, $r$	0.05
Volatility, $\sigma$	0.2
$S_{\text{range}}$	$[0, 160]$
$t_{\text{range}}$	$[0, 1]$

### 4.1 SPX Spot vs. European Call Option Price

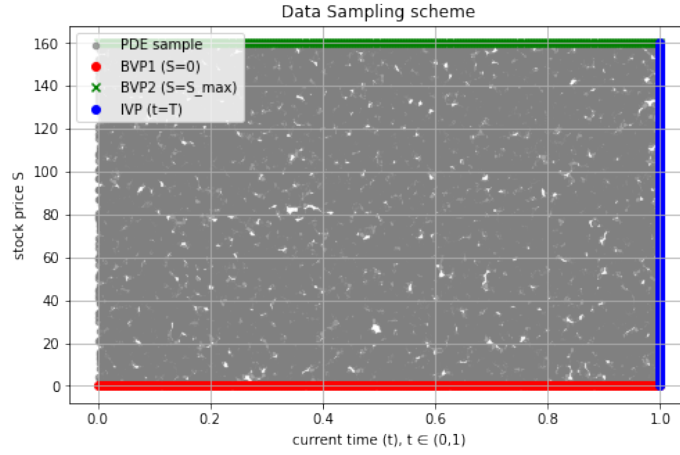
Figure 1 compares a simulated SPX spot price trajectory against the European call option price over a normalized time horizon  $t \in [0, 1]$ , using the parameters in Table 1. The SPX price (blue line) fluctuates, while the European call option price (red line) follows a smoother and consistently lower curve. This pattern aligns with financial intuition: the call price is always bounded by the underlying asset price. As  $S$  increases, the call price rises but remains tempered by the discounted strike and residual time to maturity. As  $S$  decreases, the call price approaches zero, reflecting a diminished probability of an in-the-money outcome at expiration.

### 4.2 Data Sampling Scheme

Figure 2 illustrates the data sampling scheme used for training the PINN. Gray points represent PDE collocation samples placed throughout the  $(S, t)$  domain. Red points at  $S = 0$  (BVP1) and green points at  $S = S_{\text{max}} = 160$  (BVP2) enforce boundary conditions that capture the correct behavior as  $S \rightarrow 0$  and  $S \rightarrow S_{\text{max}}$ . Blue points at  $t = T = 1$  ensure the terminal condition  $V(S, T) = \max(S - K, 0)$ . By including these sets of points, the PINN is guided to satisfy the PDE, boundary conditions, and terminal conditions simultaneously.



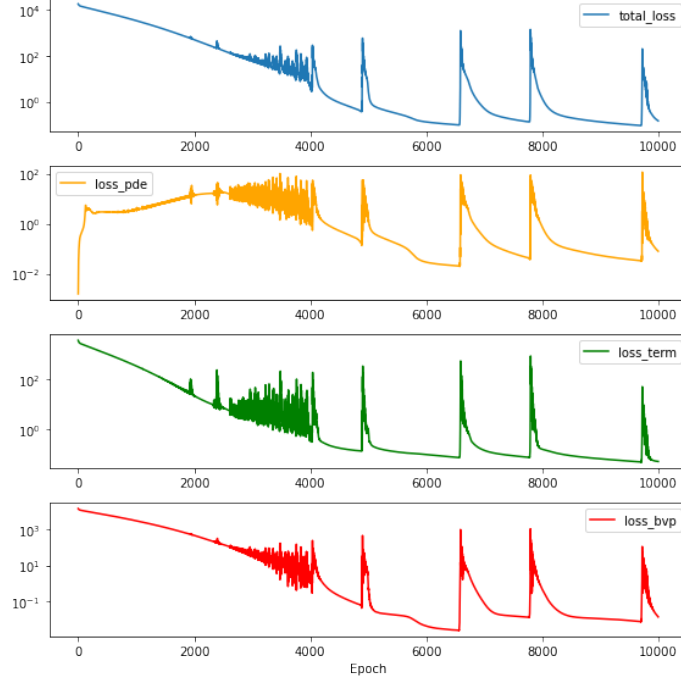
**Fig. 1** SPX spot price vs. European call option price as a function of normalized time. The blue line represents the spot price, and the red line shows the corresponding European call price.



**Fig. 2** Data sampling scheme for the PINN training process. Gray points represent PDE collocation samples, red points at  $S = 0$  and green points at  $S = S_{\max}$  enforce boundary conditions, and blue points at  $t = T$  ensure the correct terminal payoff.

### 4.3 Training Convergence and Loss Curves

Figure 3 displays the total loss and its components (PDE, terminal, and boundary losses) over training epochs. Initially, all losses decrease rapidly as the network begins to approximate the PDE solution and adhere to the imposed conditions. While some fluctuations occur, the general trend shows a downward trajectory, indicating that the PINN converges to a stable, accurate approximation. The terminal and boundary losses eventually stabilize near low values, confirming that the final network solution respects the known terminal payoff and boundary behaviors.



**Fig. 3** Training loss curves showing the total loss and individual components (PDE, terminal, and boundary) over the training epochs. The general trend indicates a decreasing loss, suggesting convergence to a stable, accurate solution.

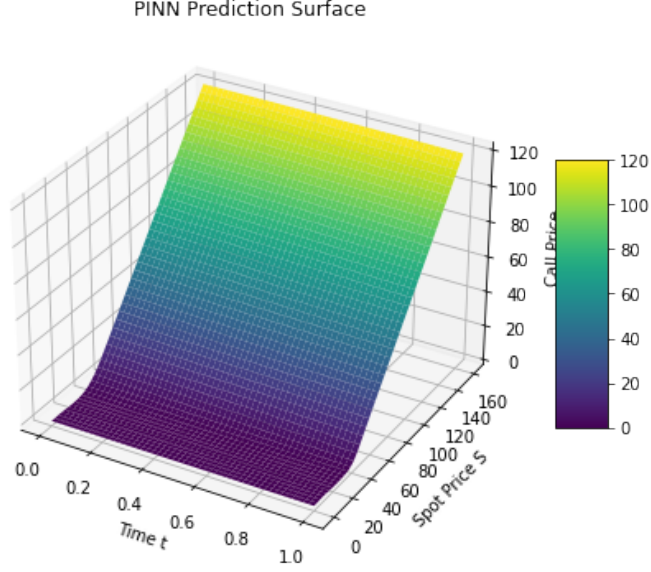
#### 4.4 PINN Predictions for the European Call Option Surface

Figure 4 shows a 3D visualization of the PINN-predicted option value across  $S \in [0, 160]$  and  $t \in [0, 1]$ . The resulting surface smoothly transitions from near-zero values at low  $S$  and early times to higher values as  $S$  increases and  $t$  approaches  $T$ . At  $t = 1$ , the PINN prediction closely matches the analytical solution  $\max(S - K, 0)$ , confirming that the network accurately reproduces the terminal payoff. This high-fidelity approximation across the entire domain further validates the effectiveness of the PINN approach.

### 5 Discussion

With the chosen parameters ( $K = 40$ ,  $r = 0.05$ ,  $\sigma = 0.2$ ,  $S \in [0, 160]$ , and  $t \in [0, 1]$ ) as listed in Table 1, our PINN model was trained using a carefully selected data sampling strategy. Approximately 20,000 interior points were used to enforce the PDE, complemented by points along the boundaries and at  $t = 1$  to ensure compliance with the terminal and boundary conditions. This balanced sampling allowed the PINN to gradually reduce all error components as training proceeded.

The performance of the model is evident in the figures. In Figure 1, as the simulated SPX price oscillates between roughly 4000 and 4200, the predicted call option



**Fig. 4** PINN prediction results. Top: The predicted 3D surface of the option price as a function of  $S$  and  $t$ . Bottom: The PINN’s prediction at  $t = 1$ , closely matching the analytical terminal payoff  $\max(S - K, 0)$ .

price consistently remains below the underlying asset’s price, reflecting realistic market behavior. Meanwhile, Figure 2 visually confirms that the chosen points cover the entire domain, indicating that the PINN was exposed to a diverse set of conditions to guide it toward the correct solution.

As seen in Figure 3, the training loss and its components declined steadily over 20,000 training steps. While the figure does not provide exact numerical values, the downward trend clearly indicates that the network adapted to satisfy both the PDE and the imposed conditions. By the final stages of training, the model closely matched the known payoff and boundary expectations, demonstrating stability and accuracy in the learning process.

Finally, the surface depicted in Figure 4 shows that the PINN accurately captured the option’s time and price dynamics. Near maturity ( $t = 1$ ), the network’s output almost perfectly aligns with the theoretical payoff line  $\max(S - K, 0)$  for  $S$  in  $[0, 160]$ . For earlier times ( $t < 1$ ), the resulting surface remains smooth and aligns well with what we expect from the Black-Scholes framework. This confirms that the model did not merely learn isolated points but rather a coherent solution throughout the domain.

In summary, these results attest that our PINN approach can solve the Black-Scholes PDE for a European call option under the given setup. The method avoided the complexity of traditional grid-based methods and still achieved a stable, accurate approximation. Future work may explore varying parameters or more complex scenarios, but the current findings demonstrate that PINNs are a viable and efficient tool for option pricing tasks.



## References

- [1] Song, L., Tan, Y., Yu, F., Luo, Y., Zheng, J.: Optimal approximations for the free boundary problems of the space-time fractional Black-Scholes equations using a combined physics-informed neural network. *Scientific Reports* **14** (2024)
- [2] Dhiman, A., Hu, Y.: Physics Informed Neural Network for Option Pricing. *arXiv preprint arXiv:2312.06711* (2023)